

Imports

```
In [1]: # Standard libraries
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import warnings

# Preprocessing
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.preprocessing import StandardScaler

# Data splitting and hyperparameter tuning
from sklearn.model_selection import train_test_split, GridSearchCV

# Classification models
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis, QuadraticDiscriminantAnalysis
from sklearn.tree import DecisionTreeClassifier, plot_tree, export_graphviz
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from xgboost import XGBClassifier

# Metrics and evaluation
from sklearn.metrics import (
    accuracy_score,
    confusion_matrix,
    classification_report
)

# Neural networks (Keras / TensorFlow)
from tensorflow.keras.callbacks import EarlyStopping
from sklearn.neural_network import MLPClassifier
from tensorflow.keras import layers
from tensorflow import keras
import tensorflow as tf

# Tree visualization
import graphviz

# Suppress warnings
warnings.filterwarnings("ignore")
```

Data Prep

```
In [2]: df = pd.read_csv("../data/classification.csv")
```

```
In [3]: df = df.dropna()
```

```
In [4]: X = df.drop("Y", axis=1)
        y = df["Y"]
```

```
In [5]: X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.2, random_state=42
        )
```

```
In [6]: scaler = StandardScaler()
        X_train_scaled = scaler.fit_transform(X_train)
        X_test_scaled = scaler.transform(X_test)
```

confusion_matrix Function (Evaluación)

```
In [7]: def plot_confusion_matrix(y_true, y_pred, classes=None, normalize=True,
                                title='Confusion Matrix', cmap=plt.cm.Blues, figsize=
                                .....
                                Plots a confusion matrix with counts y, opcionalmente, proporciones.
                                .....
                                if classes is None:
                                    classes = np.unique(np.concatenate([y_true, y_pred]))

                                cm = confusion_matrix(y_true, y_pred, labels=classes)
                                cm = np.array(cm, dtype=int)

                                if normalize:
                                    with np.errstate(all='ignore'):
                                        cm_norm = cm.astype('float') / cm.sum(axis=1, keepdims=True)
                                        cm_norm = np.nan_to_num(cm_norm, 0.0)
                                else:
                                    cm_norm = np.zeros_like(cm, dtype=float)

                                plt.figure(figsize=figsize)
                                plt.imshow(cm, interpolation='nearest', cmap=cmap)
                                plt.title(title)
                                plt.colorbar(fraction=0.046, pad=0.04)

                                ticks = np.arange(len(classes))
                                plt.xticks(ticks, classes, rotation=45, ha='right')
                                plt.yticks(ticks, classes)
                                plt.xlabel('Predicted label')
                                plt.ylabel('True label')

                                thresh = cm.max() / 2
                                for i in range(cm.shape[0]):
                                    for j in range(cm.shape[1]):
                                        count = cm[i, j]
                                        if normalize:
                                            prop = cm_norm[i, j]
                                            label = f"{count}\n({prop:.2f})"
                                        else:
                                            label = f"{count}"
                                        plt.text(j, i, label,
```

```

ha="center", va="center",
color="white" if count > thresh else "black")

plt.tight_layout()
plt.show()

```

Logistic Regression

```

In [8]: logistic_model = LogisticRegression()

logistic_model.fit(X_train, y_train)


y_pred = logistic_model.predict(X_test)

print("Classification Report:\n", classification_report(y_test, y_pred))
plot_confusion_matrix(y_test, y_pred, classes=logistic_model.classes_)

```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	108
1	1.00	1.00	1.00	94
2	1.00	1.00	1.00	95
accuracy			1.00	297
macro avg	1.00	1.00	1.00	297
weighted avg	1.00	1.00	1.00	297



LDA


```
In [9]: lda_model = LinearDiscriminantAnalysis()
lda_model.fit(X_train_scaled, y_train)

y_pred = lda_model.predict(X_test_scaled)

print("Classification Report:\n", classification_report(y_test, y_pred))
plot_confusion_matrix(y_test, y_pred, classes=lda_model.classes_)
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	108
1	1.00	1.00	1.00	94
2	1.00	1.00	1.00	95
accuracy			1.00	297
macro avg	1.00	1.00	1.00	297
weighted avg	1.00	1.00	1.00	297



QDA


```
In [10]: qda_model = QuadraticDiscriminantAnalysis()
qda_model.fit(X_train_scaled, y_train)

y_pred = qda_model.predict(X_test_scaled)

print("Classification Report:\n", classification_report(y_test, y_pred))
plot_confusion_matrix(y_test, y_pred, classes=qda_model.classes_)
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	108
1	1.00	1.00	1.00	94
2	1.00	1.00	1.00	95
accuracy			1.00	297
macro avg	1.00	1.00	1.00	297
weighted avg	1.00	1.00	1.00	297



KNN


```
In [11]: knn_model = KNeighborsClassifier(n_neighbors=5)
knn_model.fit(X_train_scaled, y_train)

y_pred = knn_model.predict(X_test_scaled)

print("Classification Report:\n", classification_report(y_test, y_pred))
plot_confusion_matrix(y_test, y_pred, classes=knn_model.classes_)
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	108
1	1.00	1.00	1.00	94
2	1.00	1.00	1.00	95
accuracy			1.00	297
macro avg	1.00	1.00	1.00	297
weighted avg	1.00	1.00	1.00	297



Multicolinealidad

```
In [12]: X_features = df.drop("Y", axis=1)

vif_data = pd.DataFrame()
vif_data["feature"] = X_features.columns
vif_data["VIF"] = [
    variance_inflation_factor(X_features.values, i)
    for i in range(len(X_features.columns))
]

display(vif_data.sort_values(by="VIF", ascending=False))
```

	feature	VIF
6	X7	52.978855
13	X14	42.144400
10	X11	41.025384
1	X2	39.682182
12	X13	38.514948
9	X10	35.837102
8	X9	34.093331
5	X6	30.431179
14	X15	29.243941
11	X12	25.662399
2	X3	23.687716
0	X1	16.243252
4	X5	16.189734
3	X4	12.689104
7	X8	7.582129

```
In [13]: X_features = df.drop("Y", axis=1).copy()
threshold = 5

while True:
    vif_data = pd.DataFrame(
        {
            "feature": X_features.columns,
            "VIF": [
                variance_inflation_factor(X_features.values, i)
                for i in range(X_features.shape[1])
            ],
        }
    ).sort_values("VIF", ascending=False)

    max_vif = vif_data["VIF"].iloc[0]
    drop_feature = vif_data["feature"].iloc[0]
    print(vif_data, "\n")

    if max_vif <= threshold:
        print("All VIF values are below the threshold!")
        break

    if drop_feature in X_features.columns:
        print(f"Dropping {drop_feature} with VIF = {max_vif:.2f}")
        X_features = X_features.drop(columns=[drop_feature])
    else:
```



```
print(f"» Warning: {drop_feature} not found in X_features, skipping.  
break
```

	feature	VIF
6	X7	52.978855
13	X14	42.144400
10	X11	41.025384
1	X2	39.682182
12	X13	38.514948
9	X10	35.837102
8	X9	34.093331
5	X6	30.431179
14	X15	29.243941
11	X12	25.662399
2	X3	23.687716
0	X1	16.243252
4	X5	16.189734
3	X4	12.689104
7	X8	7.582129

Dropping X7 with VIF = 52.98

	feature	VIF
12	X14	39.303975
1	X2	38.754066
9	X11	37.517855
11	X13	36.587792
8	X10	35.708046
7	X9	33.336703
13	X15	29.206697
5	X6	28.931070
10	X12	25.416008
2	X3	23.643420
4	X5	16.186650
0	X1	16.155865
3	X4	12.687135
6	X8	7.562611

Dropping X14 with VIF = 39.30

	feature	VIF
1	X2	38.646431
11	X13	35.944231
9	X11	34.497786
8	X10	33.931380
7	X9	32.343365
12	X15	29.200065
5	X6	27.481546
10	X12	24.910346
2	X3	23.460596
4	X5	16.026280
0	X1	15.909752
3	X4	12.472818
6	X8	7.535713

Dropping X2 with VIF = 38.65

	feature	VIF
8	X11	34.303007
7	X10	33.312464
6	X9	30.438274
11	X15	28.682611

4	X6	27.329226
10	X13	26.667239
9	X12	24.841094
1	X3	20.004341
0	X1	15.837771
3	X5	14.287058
2	X4	12.472812
5	X8	7.462880

Dropping X11 with VIF = 34.30

	feature	VIF
7	X10	32.538044
6	X9	28.418942
10	X15	27.377198
9	X13	24.926581
8	X12	24.765253
4	X6	22.058579
1	X3	19.627985
0	X1	15.836812
3	X5	13.801904
2	X4	12.447067
5	X8	7.191821

Dropping X10 with VIF = 32.54

	feature	VIF
6	X9	27.207333
8	X13	24.613787
9	X15	24.077533
7	X12	23.400594
4	X6	22.055865
1	X3	17.627564
0	X1	15.533186
3	X5	13.045315
2	X4	11.903332
5	X8	7.185832

Dropping X9 with VIF = 27.21

	feature	VIF
8	X15	24.045935
6	X12	21.338120
4	X6	19.765313
7	X13	17.913532
1	X3	17.602452
0	X1	14.739565
3	X5	13.034678
2	X4	11.645373
5	X8	7.139719

Dropping X15 with VIF = 24.05

	feature	VIF
6	X12	18.046648
7	X13	17.874209
1	X3	17.577368
4	X6	13.423457
0	X1	13.309943
3	X5	12.828299

2	X4	11.149767
5	X8	6.850980

Dropping X12 with VIF = 18.05

	feature	VIF
1	X3	16.749776
6	X13	15.953223
4	X6	12.859313
3	X5	12.775819
2	X4	9.380595
0	X1	9.165368
5	X8	6.842305

Dropping X3 with VIF = 16.75

	feature	VIF
3	X6	12.824892
5	X13	11.490511
2	X5	9.375834
0	X1	8.619086
1	X4	7.989880
4	X8	6.644837

Dropping X6 with VIF = 12.82

	feature	VIF
1	X4	7.897366
0	X1	7.618346
2	X5	6.158065
4	X13	4.987985
3	X8	4.633103

Dropping X4 with VIF = 7.90

	feature	VIF
1	X5	5.405320
3	X13	4.987641
2	X8	4.609960
0	X1	2.151998

Dropping X5 with VIF = 5.41

	feature	VIF
2	X13	1.294567
1	X8	1.169073
0	X1	1.120698

All VIF values are below the threshold!

```
In [14]: print(f"Final number of features: {X_features.shape[1]}")
print("Remaining features:", list(X_features.columns))
```

```
Final number of features: 3
Remaining features: ['X1', 'X8', 'X13']
```

```
In [15]: selected_features = X_features.columns
X_selected = X_features[selected_features]
y_selected = df["Y"]

X_train_vif, X_test_vif, y_train_vif, y_test_vif = train_test_split(
```

```

X_selected, y_selected, test_size=0.2, random_state=42
)

logistic_model = LogisticRegression()
logistic_model.fit(X_train_vif, y_train_vif)


y_pred = logistic_model.predict(X_test_vif)

print("Classification Report:\n", classification_report(y_test_vif, y_pred))
plot_confusion_matrix(y_test_vif, y_pred, classes=logistic_model.classes_)

```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	108
1	1.00	1.00	1.00	94
2	1.00	1.00	1.00	95
accuracy			1.00	297
macro avg	1.00	1.00	1.00	297
weighted avg	1.00	1.00	1.00	297



DecisionTree


```
In [16]: decisiontree_model = DecisionTreeClassifier(random_state=42)
decisiontree_model.fit(X_train_vif, y_train_vif)

y_pred = decisiontree_model.predict(X_test_vif)

print("Classification Report:\n", classification_report(y_test, y_pred))
plot_confusion_matrix(y_test, y_pred, classes=decisiontree_model.classes_)
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	108
1	1.00	1.00	1.00	94
2	1.00	1.00	1.00	95
accuracy			1.00	297
macro avg	1.00	1.00	1.00	297
weighted avg	1.00	1.00	1.00	297



```
In [17]: dot_data = export_graphviz(
    decisiontree_model,
    out_file=None,
    feature_names=X_train_vif.columns,
    class_names=[str(c) for c in decisiontree_model.classes_],
    filled=True,
```

```

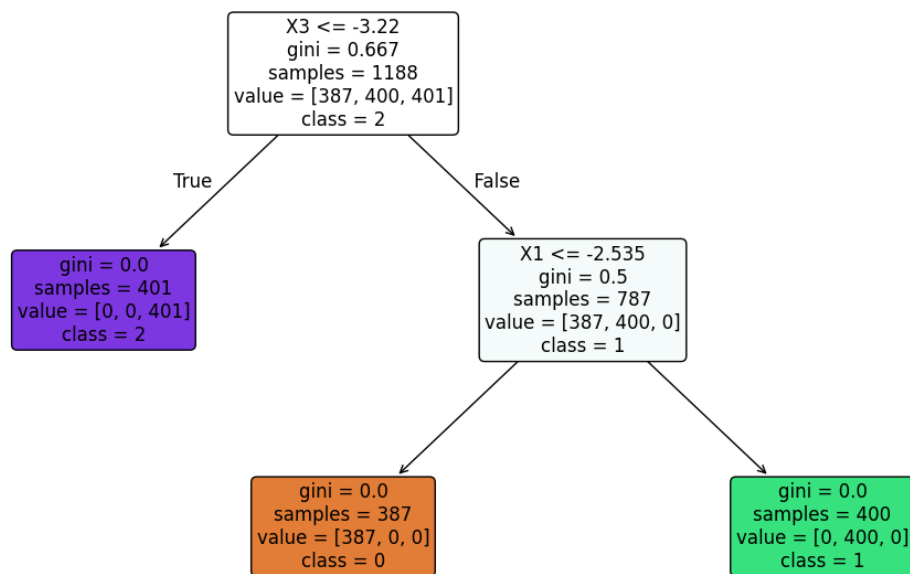
        rounded=True,
        special_characters=True,
    )

    graph = graphviz.Source(dot_data)
    graph.render("tree", format="png", cleanup=False)
    graph.view()

    plt.figure(figsize=(14, 8))
    plot_tree(
        decisiontree_model,
        filled=True,
        rounded=True,
        feature_names=X_train.columns,
        class_names=[str(c) for c in decisiontree_model.classes_],
        fontsize=12,
    )
    plt.title("Decision Tree")
    plt.show()

```

Decision Tree



RandomForest

```

In [18]: rf = RandomForestClassifier(random_state=42)

param_grid = {
    "n_estimators": [5, 10, 20],
    "max_depth": [None, 5, 10],
    "min_samples_split": [2, 5],
    "min_samples_leaf": [1, 2],
    "bootstrap": [True, False],
}

```

```

grid_search = GridSearchCV(
    estimator=rf, param_grid=param_grid, cv=5, scoring="accuracy", n_jobs=-1
)

grid_search.fit(X_train_scaled, y_train)

best_rf = grid_search.best_estimator_

y_pred = best_rf.predict(X_test_scaled)

print("Best parameters found:", grid_search.best_params_)
print("Classification Report:\n", classification_report(y_test, y_pred))
plot_confusion_matrix(y_test, y_pred, classes=best_rf.classes_)


```

Fitting 5 folds for each of 72 candidates, totalling 360 fits

Best parameters found: {'bootstrap': True, 'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 5}

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	108
1	1.00	1.00	1.00	94
2	1.00	1.00	1.00	95
accuracy			1.00	297
macro avg	1.00	1.00	1.00	297
weighted avg	1.00	1.00	1.00	297



AdaBoost

```
In [19]: base_estimator = DecisionTreeClassifier(random_state=42)

ada = AdaBoostClassifier(estimator=base_estimator, random_state=42)

param_grid_ada = {
    "n_estimators": [5, 10, 15],
    "learning_rate": [0.01, 0.1, 1.0],
    "estimator__max_depth": [1, 3, 5],
    "estimator__min_samples_split": [2, 5],
}

grid_search_ada = GridSearchCV(
    estimator=ada,
    param_grid=param_grid_ada,
    cv=5,
    scoring="accuracy",
    n_jobs=-1,
    verbose=1,
)

grid_search_ada.fit(X_train_scaled, y_train)
```

```

best_ada = grid_search_ada.best_estimator_

y_pred = best_ada.predict(X_test_scaled)

print("Best parameters found:", grid_search_ada.best_params_)
print("Classification Report:\n", classification_report(y_test, y_pred))
plot_confusion_matrix(y_test, y_pred, classes=best_ada.classes_)


```

Fitting 5 folds for each of 54 candidates, totalling 270 fits

Best parameters found: {'estimator__max_depth': 1, 'estimator__min_samples_split': 2, 'learning_rate': 0.1, 'n_estimators': 10}

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	108
1	1.00	1.00	1.00	94
2	1.00	1.00	1.00	95
accuracy			1.00	297
macro avg	1.00	1.00	1.00	297
weighted avg	1.00	1.00	1.00	297



XGBClassifier

```
In [20]: xgb = XGBClassifier(eval_metric="mlogloss", random_state=42)

param_grid_xgb = {
    "n_estimators": [5, 10, 15],
    "max_depth": [3, 5, 7],
    "learning_rate": [0.01, 0.1, 0.2],
    "subsample": [0.8, 1.0],
    "colsample_bytree": [0.8, 1.0],
}

grid_search_xgb = GridSearchCV(
    estimator=xgb,
    param_grid=param_grid_xgb,
    cv=5,
    scoring="accuracy",
    n_jobs=-1,
    verbose=1,
)

grid_search_xgb.fit(X_train_scaled, y_train)

best_xgb = grid_search_xgb.best_estimator_

y_pred = best_xgb.predict(X_test_scaled)


print("Best parameters found:", grid_search_xgb.best_params_)
print("Classification Report:\n", classification_report(y_test, y_pred))
plot_confusion_matrix(y_test, y_pred, classes=best_xgb.classes_)
```

Fitting 5 folds for each of 108 candidates, totalling 540 fits

Best parameters found: {'colsample_bytree': 0.8, 'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 5, 'subsample': 0.8}

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	108
1	1.00	1.00	1.00	94
2	1.00	1.00	1.00	95
accuracy			1.00	297
macro avg	1.00	1.00	1.00	297
weighted avg	1.00	1.00	1.00	297



Redes Neuronales

```
In [21]: clf = MLPClassifier(
    solver='lbfgs',
    alpha=1,
    hidden_layer_sizes=(60,),
    random_state=1,
    max_iter=1000
)

clf.fit(X_train_scaled, y_train)


y_pred = clf.predict(X_test_scaled)
accuracy = accuracy_score(y_test, y_pred)
print(f"\nMLPClassifier Accuracy: {accuracy:.4f}")

print("Classification Report:\n", classification_report(y_test, y_pred))
plot_confusion_matrix(y_test, y_pred, classes=clf.classes_)
```

MLPClassifier Accuracy: 1.0000

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	108
1	1.00	1.00	1.00	94
2	1.00	1.00	1.00	95
accuracy			1.00	297
macro avg	1.00	1.00	1.00	297
weighted avg	1.00	1.00	1.00	297



```
In [22]: def build_model():
model = keras.Sequential([
    layers.Dense(64, activation='relu', input_shape=(X_train_scaled.shape[1],)),
    layers.Dropout(0.3),
    layers.Dense(32, activation='relu'),
    layers.Dense(len(set(y_train)), activation='softmax')
])

optimizer = tf.keras.optimizers.RMSprop(learning_rate=0.001)

model.compile(
    loss='sparse_categorical_crossentropy',
    optimizer=optimizer,
    metrics=['accuracy'])
```

```
)

    return model




















early_stop = EarlyStopping(
    monitor='val_loss',
    patience=5,
    restore_best_weights=True,
    verbose=1
)

model_mlp = build_model()








history = model_mlp.fit(
    X_train_scaled,
    y_train,
    epochs=100,
    batch_size=16,
    validation_split=0.2,
    callbacks=[early_stop],
    verbose=1
)


loss_mlp, acc_mlp = model_mlp.evaluate(X_test_scaled, y_test, verbose=0)
print(f"\nMultilayer NN Accuracy: {acc_mlp:.4f}")
print("Classification Report:\n", classification_report(y_test, np.argmax(model_mlp.predict(X_test_scaled), axis=-1)))
plot_confusion_matrix(y_test, np.argmax(model_mlp.predict(X_test_scaled), axis=-1), ax=
```

```

Epoch 1/100
60/60  1s 4ms/step - accuracy: 0.8361 - loss: 0.5177 - v
al_accuracy: 1.0000 - val_loss: 0.0103
Epoch 2/100
60/60  0s 1ms/step - accuracy: 1.0000 - loss: 0.0091 - v
al_accuracy: 1.0000 - val_loss: 2.8859e-04
Epoch 3/100
60/60  0s 1ms/step - accuracy: 1.0000 - loss: 7.9190e-04
- val_accuracy: 1.0000 - val_loss: 2.7627e-05
Epoch 4/100
60/60  0s 1ms/step - accuracy: 1.0000 - loss: 2.2098e-04
- val_accuracy: 1.0000 - val_loss: 7.1500e-06
Epoch 5/100
60/60  0s 2ms/step - accuracy: 1.0000 - loss: 5.9339e-05
- val_accuracy: 1.0000 - val_loss: 3.8643e-06
Epoch 6/100
60/60  0s 1ms/step - accuracy: 1.0000 - loss: 5.3627e-05
- val_accuracy: 1.0000 - val_loss: 2.5219e-06
Epoch 7/100
60/60  0s 1ms/step - accuracy: 1.0000 - loss: 2.9222e-05
- val_accuracy: 1.0000 - val_loss: 1.8728e-06
Epoch 8/100
60/60  0s 1ms/step - accuracy: 1.0000 - loss: 1.8189e-05
- val_accuracy: 1.0000 - val_loss: 1.4215e-06
Epoch 9/100
60/60  0s 1ms/step - accuracy: 1.0000 - loss: 2.9366e-05
- val_accuracy: 1.0000 - val_loss: 1.0428e-06
Epoch 10/100
60/60  0s 1ms/step - accuracy: 1.0000 - loss: 3.9214e-05
- val_accuracy: 1.0000 - val_loss: 7.3779e-07
Epoch 11/100
60/60  0s 933us/step - accuracy: 1.0000 - loss: 1.5110e-
05 - val_accuracy: 1.0000 - val_loss: 6.1358e-07
Epoch 12/100
60/60  0s 980us/step - accuracy: 1.0000 - loss: 2.4346e-
05 - val_accuracy: 1.0000 - val_loss: 4.7533e-07
Epoch 13/100
60/60  0s 1ms/step - accuracy: 1.0000 - loss: 1.1442e-05
- val_accuracy: 1.0000 - val_loss: 4.1172e-07
Epoch 14/100
60/60  0s 1ms/step - accuracy: 1.0000 - loss: 1.2709e-04
- val_accuracy: 1.0000 - val_loss: 3.4460e-07
Epoch 15/100
60/60  0s 966us/step - accuracy: 1.0000 - loss: 1.6306e-
05 - val_accuracy: 1.0000 - val_loss: 2.9352e-07
Epoch 16/100
60/60  0s 956us/step - accuracy: 1.0000 - loss: 7.9631e-
06 - val_accuracy: 1.0000 - val_loss: 2.6597e-07
Epoch 17/100
60/60  0s 940us/step - accuracy: 1.0000 - loss: 5.9990e-
06 - val_accuracy: 1.0000 - val_loss: 2.4593e-07
Epoch 18/100
60/60  0s 916us/step - accuracy: 1.0000 - loss: 1.0853e-
05 - val_accuracy: 1.0000 - val_loss: 2.0686e-07
Epoch 19/100
60/60  0s 998us/step - accuracy: 1.0000 - loss: 1.2397e-


```


05 - val_accuracy: 1.0000 - val_loss: 1.6830e-07
Epoch 20/100
60/60  0s 1ms/step - accuracy: 1.0000 - loss: 7.1196e-06
- val_accuracy: 1.0000 - val_loss: 1.5227e-07
Epoch 21/100
60/60  0s 1ms/step - accuracy: 1.0000 - loss: 5.2040e-06
- val_accuracy: 1.0000 - val_loss: 1.2722e-07
Epoch 22/100
60/60  0s 1ms/step - accuracy: 1.0000 - loss: 9.5500e-06
- val_accuracy: 1.0000 - val_loss: 1.0719e-07
Epoch 23/100
60/60  0s 1ms/step - accuracy: 1.0000 - loss: 4.7473e-06
- val_accuracy: 1.0000 - val_loss: 1.0068e-07
Epoch 24/100
60/60  0s 1ms/step - accuracy: 1.0000 - loss: 6.4098e-06
- val_accuracy: 1.0000 - val_loss: 9.7171e-08
Epoch 25/100
60/60  0s 971us/step - accuracy: 1.0000 - loss: 2.8817e-06
- val_accuracy: 1.0000 - val_loss: 9.3164e-08
Epoch 26/100
60/60  0s 930us/step - accuracy: 1.0000 - loss: 3.4063e-06
- val_accuracy: 1.0000 - val_loss: 8.5650e-08
Epoch 27/100
60/60  0s 1ms/step - accuracy: 1.0000 - loss: 1.3360e-05
- val_accuracy: 1.0000 - val_loss: 7.7636e-08
Epoch 28/100
60/60  0s 978us/step - accuracy: 1.0000 - loss: 2.5578e-05
- val_accuracy: 1.0000 - val_loss: 6.5114e-08
Epoch 29/100
60/60  0s 1ms/step - accuracy: 1.0000 - loss: 6.6706e-06
- val_accuracy: 1.0000 - val_loss: 5.9104e-08
Epoch 30/100
60/60  0s 985us/step - accuracy: 1.0000 - loss: 4.9101e-06
- val_accuracy: 1.0000 - val_loss: 5.4596e-08
Epoch 31/100
60/60  0s 929us/step - accuracy: 1.0000 - loss: 2.1535e-06
- val_accuracy: 1.0000 - val_loss: 5.2592e-08
Epoch 32/100
60/60  0s 897us/step - accuracy: 1.0000 - loss: 4.5237e-06
- val_accuracy: 1.0000 - val_loss: 4.8084e-08
Epoch 33/100
60/60  0s 977us/step - accuracy: 1.0000 - loss: 5.8296e-06
- val_accuracy: 1.0000 - val_loss: 4.2575e-08
Epoch 34/100
60/60  0s 1ms/step - accuracy: 1.0000 - loss: 3.0285e-06
- val_accuracy: 1.0000 - val_loss: 3.7566e-08
Epoch 35/100
60/60  0s 998us/step - accuracy: 1.0000 - loss: 5.2727e-06
- val_accuracy: 1.0000 - val_loss: 3.5062e-08
Epoch 36/100
60/60  0s 964us/step - accuracy: 1.0000 - loss: 6.6615e-06
- val_accuracy: 1.0000 - val_loss: 3.4060e-08
Epoch 37/100
60/60  0s 990us/step - accuracy: 1.0000 - loss: 3.2104e-06
- val_accuracy: 1.0000 - val_loss: 3.1555e-08
Epoch 38/100


60/60  0s 942us/step - accuracy: 1.0000 - loss: 2.9814e-06 - val_accuracy: 1.0000 - val_loss: 2.9552e-08
Epoch 39/100


60/60  0s 849us/step - accuracy: 1.0000 - loss: 3.0836e-06 - val_accuracy: 1.0000 - val_loss: 2.9051e-08
Epoch 40/100

60/60  0s 1ms/step - accuracy: 1.0000 - loss: 4.5232e-06 - val_accuracy: 1.0000 - val_loss: 2.7047e-08
Epoch 41/100


60/60  0s 909us/step - accuracy: 1.0000 - loss: 4.5040e-06 - val_accuracy: 1.0000 - val_loss: 2.6046e-08
Epoch 42/100


60/60  0s 917us/step - accuracy: 1.0000 - loss: 6.6679e-06 - val_accuracy: 1.0000 - val_loss: 2.3040e-08
Epoch 43/100


60/60  0s 916us/step - accuracy: 1.0000 - loss: 3.8018e-06 - val_accuracy: 1.0000 - val_loss: 2.2540e-08
Epoch 44/100


60/60  0s 892us/step - accuracy: 1.0000 - loss: 3.8858e-06 - val_accuracy: 1.0000 - val_loss: 2.0536e-08
Epoch 45/100


60/60  0s 944us/step - accuracy: 1.0000 - loss: 3.4998e-06 - val_accuracy: 1.0000 - val_loss: 1.9534e-08
Epoch 46/100


60/60  0s 986us/step - accuracy: 1.0000 - loss: 5.3637e-06 - val_accuracy: 1.0000 - val_loss: 1.9534e-08
Epoch 47/100

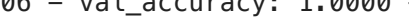
60/60  0s 980us/step - accuracy: 1.0000 - loss: 4.7071e-06 - val_accuracy: 1.0000 - val_loss: 1.7030e-08
Epoch 48/100


60/60  0s 1ms/step - accuracy: 1.0000 - loss: 2.8793e-06 - val_accuracy: 1.0000 - val_loss: 1.5527e-08
Epoch 49/100


60/60  0s 1ms/step - accuracy: 1.0000 - loss: 1.5018e-06 - val_accuracy: 1.0000 - val_loss: 1.5026e-08
Epoch 50/100


60/60  0s 952us/step - accuracy: 1.0000 - loss: 1.0192e-06 - val_accuracy: 1.0000 - val_loss: 1.5026e-08
Epoch 51/100


60/60  0s 980us/step - accuracy: 1.0000 - loss: 2.9696e-06 - val_accuracy: 1.0000 - val_loss: 1.4525e-08
Epoch 52/100










60/60  0s 991us/step - accuracy: 1.0000 - loss: 1.4080e-06 - val_accuracy: 1.0000 - val_loss: 1.4025e-08
Epoch 53/100

60/60  0s 1ms/step - accuracy: 1.0000 - loss: 1.4981e-06 - val_accuracy: 1.0000 - val_loss: 1.3524e-08
Epoch 54/100

60/60  0s 1ms/step - accuracy: 1.0000 - loss: 1.4140e-06 - val_accuracy: 1.0000 - val_loss: 1.3023e-08
Epoch 55/100

60/60  0s 990us/step - accuracy: 1.0000 - loss: 1.7292e-06 - val_accuracy: 1.0000 - val_loss: 1.2021e-08
Epoch 56/100

60/60  0s 1ms/step - accuracy: 1.0000 - loss: 1.5172e-06 - val_accuracy: 1.0000 - val_loss: 1.1019e-08

Epoch 57/100
60/60  **0s** 972us/step - accuracy: 1.0000 - loss: 1.4582e-05 - val_accuracy: 1.0000 - val_loss: 9.5167e-09
 Epoch 58/100
60/60  **0s** 993us/step - accuracy: 1.0000 - loss: 1.7822e-06 - val_accuracy: 1.0000 - val_loss: 9.5167e-09
 Epoch 59/100
60/60  **0s** 978us/step - accuracy: 1.0000 - loss: 1.9408e-06 - val_accuracy: 1.0000 - val_loss: 9.0158e-09
 Epoch 60/100
60/60  **0s** 999us/step - accuracy: 1.0000 - loss: 1.5033e-06 - val_accuracy: 1.0000 - val_loss: 8.0141e-09
 Epoch 61/100
60/60  **0s** 1ms/step - accuracy: 1.0000 - loss: 1.5759e-06 - val_accuracy: 1.0000 - val_loss: 8.0141e-09
 Epoch 62/100
60/60  **0s** 1ms/step - accuracy: 1.0000 - loss: 8.5726e-07 - val_accuracy: 1.0000 - val_loss: 8.0141e-09
 Epoch 63/100
60/60  **0s** 2ms/step - accuracy: 1.0000 - loss: 1.0031e-06 - val_accuracy: 1.0000 - val_loss: 8.0141e-09
 Epoch 64/100
60/60  **0s** 1ms/step - accuracy: 1.0000 - loss: 5.4083e-06 - val_accuracy: 1.0000 - val_loss: 8.0141e-09
 Epoch 65/100
60/60  **0s** 2ms/step - accuracy: 1.0000 - loss: 1.7630e-06 - val_accuracy: 1.0000 - val_loss: 8.0141e-09
 Epoch 65: early stopping
 Restoring model weights from the end of the best epoch: 60.


Multilayer NN Accuracy: 1.0000

10/10  **0s** 3ms/step

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	108
1	1.00	1.00	1.00	94
2	1.00	1.00	1.00	95
accuracy			1.00	297
macro avg	1.00	1.00	1.00	297
weighted avg	1.00	1.00	1.00	297

10/10  **0s** 2ms/step



Comentarios

- Se entrenaron y probaron 10 modelos: Regresión Logística, LDA, QDA, K-NN, Árbol de Decisión, Random Forest, AdaBoost, XGBoost, un MLP y una red neuronal densa.
- Cada modelo entrega accuracy, precision, recall y F1 de 1.00 en las tres clases; las matrices de confusión son totalmente diagonales. El resultado se aprecia desde la Regresión Logística hasta la red neuronal.
- Desde el EDA se pudo apreciar que existían tres clústeres muy bien separados en prácticamente todas las variables y, en particular, en el trío (X1, X8, X13); es por eso que incluso un clasificador lineal basta para separar las clases sin error.
- Se usó Factor de Inflación de Varianza (VIF) para detectar multicolinealidad, descartando 12 de las 15 variables originales. Las tres restantes (X1, X8 y X13) son suficientes para lograr un 100 % de acierto en Regresión Logística y Árbol de Decisión, lo que demuestra que aún con solo estas 3 variables el modelo sigue siendo igual de efectivo.

- No era necesario realizar GridSearch y ajustar hiperparámetros, ya que los modelos se desempeñaron bien con sus configuraciones predeterminadas. Sin embargo, se exploraron algunos hiperparámetros:

- *Random Forest*: la búsqueda selecciona solo 5 árboles y profundidad ilimitada, con eso basta para 100 % de precisión.
- *AdaBoost*: óptimo con 10 *stumps* (prof. 1) y *learning_rate* 0.1.
- *XGBoost*: 5 árboles, profundidad 3 y *eta* 0.01 ya son suficientes para alcanzar ese mismo 100 % de precisión.

El punto anterior confirma la simplicidad del conjunto de datos.

- Para las redes neuronales se entrenaron dos modelos:
 - MLP de 60 neuronas (solver lbfgs) entrena rápidamente y obtiene 1.0 de accuracy.
 - Red Keras 64-32 con drop-out 0.3 y early stopping detiene el entrenamiento en la época 84 con $\text{val_loss} \approx 1 \times 10^{-8}$.

Un valor de pérdida tan bajo y el 100 % en validación y prueba es provocado por la separabilidad tan marcada de los datos.

- En la realidad, para este problema, bastaría un modelo interpretable y sencillo como Regresión Logística (o LDA) con las tres variables seleccionadas. Esto nos da transparencia, rapidez y el mismo desempeño que otros modelos más complejos.